

What is Text Classification

Text classification is automatically classifying a document into two or more *predefined* categories.

Examples of documents and their possible categories: Emails: spam or not spam Newswire articles: business, politics or sports
Movie reviews: liked it or hated it

Supervised learning

What knowledge do we need to complete the task?

Classifying newswires into politics or sports: Read the newswires, find the characteristics which determine whether they are politics or sports, then make a judgement.

But what if you've never read a politics article before? How would you know what to look for?

Supervised learning

What knowledge do we need to complete the task?

Classifying newswires into politics or sports: Read the newswires, find the characteristics which determine whether they are politics or sports, then make a judgement.

But what if you've never read a politics article before? How would you know what to look for?

In *supervised learning*, the machine learns based on examples and their **already known** classifications (*training set*)

Based on the knowledge learnt from the examples, we can classify **unknown** examples (*test set*)

Today's problem

Lady Gaga vs. The Clash

Input to this problem: a bunch of text files containing song lyrics.

Preprocessing the input

Decompose (tokenise) each input file into a 'bag of words'.

Bag of words: a list of words with corresponding frequency of that word occurring in the document (a weighted vector of words)

	<i>word₁</i>	<i>word₂</i>	<i>word₃</i>	<i>word₄</i>	<i>word₅</i>
<i>doc₁</i>					
<i>doc₂</i>					
<i>doc₃</i>					

Train the classifier

Multinomial Naive Bayes: a Generative learning algorithm

Generative learning algorithms create a model based on the training data, and apply that model on the test data.

Two categories: *gaga*, *clash*

We will build a model for each category.

Classification task: find the *most likely* category that the document belongs to.

Demo

How does it work?

Bayes' Theorem

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Where:

A and B are events

$P(B) \neq 0$ and is the probability of the data

$P(A|B)$: probability that A occurs given B

$P(B|A)$: probability that B occurs given A

Bayes' Theorem in Text Classification

Objective: find the likelihood that document d belongs in class c

$$P(c|d) = \frac{P(d|c)P(c)}{P(d)}$$

Try each of the classes in turn i.e. $C = \{gaga, clash\}$

The most likely class is the one which returns the highest value of $P(c|d)$ which is defined as C_{MAP} (maximum a posteriori)

$$C_{MAP} = \arg \max_{c \in C} P(c|d)$$

Bayes' Theorem in Text Classification

Objective: find the likelihood that document d belongs in class c

$$P(c|d) = \frac{P(d|c)P(c)}{P(d)}$$

Try each of the classes in turn i.e. $C = \{gaga, clash\}$

The most likely class is the one which returns the highest value of $P(c|d)$ which is defined as C_{MAP} (maximum a posteriori)

$$C_{MAP} = \arg \max_{c \in C} P(c|d)$$

(Apply Bayes' theorem again)

$$C_{MAP} = \arg \max_{c \in C} \frac{P(d|c)P(c)}{P(d)}$$

$$C_{MAP} = \arg \max_{c \in C} P(d|c)P(c)$$

Training

Recall that our document (d) is just a 'bag of words'. We can represent the words as the product of all of the word probabilities:

$$P(c|d) \propto P(c) \prod_{1 \leq k \leq n_d} P(t_k|c)$$

Prior probability of the class:

$$P(C) = \frac{N_c}{N}$$

Where N_c is the number of docs in the class, and N is the total number of docs.

Relative frequency of term t occurring in a class c :

$$P(t|c) = \frac{T_{ct}}{\sum_{t \in V} T_{ct}}$$

where T_{ct} is the number of occurrences of t in training documents from class c . (V is the set of all terms)

The problem of zero

Since we need to iterate through all terms in our vocabulary for all classes, if a document in a class doesn't include a term existing in the other classes, $T_{ct} = 0$, which makes $P(c|d) = 0$ when included in the product of all terms.

For example, *gaga* documents include the word 'mascara' which appears in no *clash* documents, so $P(mascara|clash) = 0$.

We fix this by applying *add-one smoothing*:

$$P(t|c) = \frac{T_{ct}}{\sum_{t \in V} T_{ct}}$$

becomes

$$P(t|c) = \frac{T_{ct} + 1}{\sum_{t \in V} (T_{ct} + 1)} = \frac{T_{ct} + 1}{(\sum_{t \in V} T_{ct}) + |V|}$$

where $|V|$ is the total number of terms in the vocabulary.

Example: training

Training set:

	N_c	animal	game	love	london	T_{ct}
gaga	2	66	1	21	0	88
clash	2	0	4	0	14	18

Class prior: $P(c) = \frac{N_c}{N} = \frac{2}{4}$ for both classes so we can dismiss it.

Term likelihood:

	animal	game	love	london
gaga	$P(\text{animal} \text{gaga})$			
clash				

$$P(t|c) = \frac{T_{ct}+1}{(\sum_{t \in V} T_{ct}) + |V|}$$

$$P(\text{animal}|\text{gaga}) = \frac{66+1}{88+106} = 0.345361$$

Example: training

Training set:

	N_c	animal	game	love	london	T_{ct}
gaga	2	66	1	21	0	88
clash	2	0	4	0	14	18

Class prior: $P(c) = \frac{N_c}{N} = \frac{2}{4}$ for both classes so we can dismiss it.

Term likelihood:

	animal	game	love	london
gaga	0.345361			
clash	$P(\text{animal} \text{clash})$			

$$P(t|c) = \frac{T_{ct}+1}{(\sum_{t \in V} T_{ct}) + |V|}$$

$$P(\text{animal}|\text{gaga}) = \frac{66+1}{88+106} = 0.345361$$

$$P(\text{animal}|\text{clash}) = \frac{0+1}{18+106} = 0.008065$$

Example: training

Training set:

	N_c	animal	game	love	london	T_{ct}
gaga	2	66	1	21	0	88
clash	2	0	4	0	14	18

Class prior: $P(c) = \frac{N_c}{N} = \frac{2}{4}$ for both classes so we can dismiss it.

Term likelihood:

	animal	game	love	london
gaga	0.345361			
clash	0.008065			

$$P(t|c) = \frac{T_{ct}+1}{(\sum_{t \in V} T_{ct}) + |V|}$$

$$P(\text{animal}|\text{gaga}) = \frac{66+1}{88+106} = 0.345361$$

$$P(\text{animal}|\text{clash}) = \frac{0+1}{18+106} = 0.008065$$

Example: training

Training set:

	N_c	animal	game	love	london	T_{ct}
gaga	2	66	1	21	0	88
clash	2	0	4	0	14	18

Class prior: $P(c) = \frac{N_c}{N} = \frac{2}{4}$ for both classes so we can dismiss it.

Term likelihood:

	animal	game	love	london
gaga	0.345361	0.010309	0.113402	0.005155
clash	0.008065	0.040323	0.008605	0.120968

$$P(t|c) = \frac{T_{ct}+1}{(\sum_{t \in V} T_{ct}) + |V|}$$

$$P(\text{animal}|\text{gaga}) = \frac{66+1}{88+106} = 0.345361$$

$$P(\text{animal}|\text{clash}) = \frac{0+1}{18+106} = 0.008065$$

Example: classifying

Our trained classifier:

	animal	game	love	london
gaga	0.345361	0.010309	0.113402	0.005155
clash	0.008065	0.040323	0.008605	0.120968

Test document:

	animal	game	love	london
???		1	2	1

Example: classifying

Our trained classifier:

	animal	game	love	london
gaga	0.345361	0.010309	0.113402	0.005155
clash	0.008065	0.040323	0.008605	0.120968

Test document:

	animal	game	love	london
???		1	2	1

Multiply the test document by the training set:

	animal	game	love	london
gaga		0.010309	0.226804	0.005155
clash		0.040323	0.016129	0.120968

Example: classifying

Our trained classifier:

	animal	game	love	london
gaga	0.345361	0.010309	0.113402	0.005155
clash	0.008065	0.040323	0.008605	0.120968

Test document:

	animal	game	love	london
???		1	2	1

Multiply the test document by the training set:

	animal	game	love	london
gaga		0.010309	0.226804	0.005155
clash		0.040323	0.016129	0.120968

This gives:

$$P(\text{gaga}) = 0.010309 * 0.226804 * 0.005155 = 6.02625e^{-6}$$

$$P(\text{clash}) = 0.040323 * 0.016129 * 0.120968 = 3.93365e^{-5}$$

Example: classifying

Our trained classifier:

	animal	game	love	london
gaga	0.345361	0.010309	0.113402	0.005155
clash	0.008065	0.040323	0.008605	0.120968

Test document:

	animal	game	love	london
???		1	2	1

Multiply the test document by the training set:

	animal	game	love	london
gaga		0.010309	0.226804	0.005155
clash		0.040323	0.016129	0.120968

This gives:

$$P(\text{gaga}) = 0.010309 * 0.226804 * 0.005155 = 6.02625e^{-6}$$

$$P(\text{clash}) = 0.040323 * 0.016129 * 0.120968 = 3.93365e^{-5}$$

Wins

References

Christopher D. Manning, Prabhakar Raghavan and Hinrich Schtze.
(2008) *Introduction to Information Retrieval*, Cambridge University
Press. <https://nlp.stanford.edu/IR-book/>